

Task 2: how many more visitors are increased in weekends than week days

Step1: To do this job I have created hive table

```
create table ckin_pivot as select business_id ,COUNT(IF(date='mon',business_id,null)) as
monday,COUNT(IF(date='tue',business_id,null)) as tuesday,COUNT(IF(date
='wed',business_id,null)) as wednesday,COUNT(IF(date='thu',business_id,null)) as
thursday,COUNT(IF(date='fri',business_id,null))as friday,COUNT(IF(date='sat',business_id,null))
as saturday,COUNT(IF(date='sun',business_id,null))as sunday from ckin c where c.date
IN('mon'|'tue'|'wed') GROUP BY business_id;
```

BusinessID Monday. Tuesday. Wednesday Thursday Friday Saturday Sunday

```
hive> select * from ckin_pivot;
```

```
OK
-49WY_TeA9ZEcRk_GnuLog 4      0      3      0      2      1      0
05yzn45mHlTCbT25KmFVlQ 0      0      0      0      2      9      0
0dcNFXs5yMjln2RtkgIoHg 4      4      3      1      1      5      1
0FL1-ZegpmlG8P9kZEoZjw 0      0      1      0      0      0      0
0Fln0xZBGfxvRd66pbZ3ag 1      0      2      2      6      12     10
0VJlqjDftGcrtgRhbkRtQA 0      0      0      1      0      1      0
0wLKvB8ui5SZdlj-r8EWtQ 1      0      0      0      0      0      0
0WsCtzG8xazAmFA5HhQE0A 2      1      0      1      1      3      4
1-rUjW65plaAjIMVMjJP_Q 1      0      0      1      3      1      0
1EH3UCsWkTHH4Wlwed72JA 0      0      0      0      8      0      0
1RV1FeZKJN7JJlnZpEIVWw 0      0      1      0      0      1      0
1Xi4hGRzhXZleq37hdFLcg 0      0      2      0      0      6      0
243PQVgkx8eUzVaKCGC-gw 0      0      1      0      0      7      11
2pm0I_mrn1ZWB88IeiOMkA 0      0      0      0      0      0      4
2wfiUn6AZm2K00FJYakVQ 0      0      0      0      0      9      0
2X9U7v-AvoiB-ki0y85bA 5      5      4      2      2      3      1
2YIg-PAGCMXMLxw53T29Q 0      0      6      4      0      0      2
3XBN60tLWqr7WZuDZsyZkA 0      0      0      0      0      11     0
3ZQxw4vY0HBuS7KeiWmyew 7      3      5      3      5      6      0
49WY_TeA9ZEcRk_GnuLog 8      0      6      0      11     14     15
4iuFGkrZYnWZg0axNantdw 6      5      0      0      0      0      0
4k2o2LNXm9ibZ1lXCV80Ug 4      6      7      10     6      2      8
4z0yZff4-FemUqRW3GXuCA 3      1      1      2      2      2      1
5A2BcDVbWGM7JEUgSfFleA 0      0      0      2      1      4      5
5YyLWps6xL372Lj2hRlPQw 9      5      11     19     7      13     12
6FkCrXJMq5KIgavHRqtlbQ 0      0      1      0      0      2      3
6gDRAYnqwHDirfA5gFlsYg 0      0      0      0      0      3      0
6iq7LAJRSstbaYcLzQ2EzA 8      8      5      6      8      12     13
6PgUz2MCsI8V-7utG_42LA 0      0      1      0      1      2      0
6pLZrlG99_Ig80kbsqo28A 2      4      4      9      4      4      3
6Rxt0_MxnLwZ7D7iYDE8Jg 9      5      7      4      7      11     11
7KPBkxA0EtB3QeIL9PEErg 10     10     12     11     12     13     12
7rHrh5QgqFBXj40jF0NDCA 2      4      0      1      2      1      6
7sakpP2av046jujj-ERHhg 11     4      6      12     18     23     18
8bY6M2yiW0F2ilfmGS34Fw 0      0      0      1      1      1      0
8ZL4ed3Y8_F6AgPAYxgi2g 4      0      0      0      0      1      0
a0H7WlpoiHf8S_OQzBs8Fw 2      3      3      0      1      1      1
Al-oKob7eDmt-Te9FHuzw 0      0      0      0      0      0      1
aLnnIPzhKT7CBMACH95CQg 1      4      0      4      3      0      0
```

Step2:

create table difference as select business_id,ROUND(saturday/2+sunday/2)-
ROUND(monday/4+tuesday/4+thursday/4+wednesday/4+friday/4) as weekendmorevisitors
from ckin_pivot where ROUND(monday/4+tuesday/4+thursday/4+wednesday/4+friday/4) <
ROUND(saturday/2+sunday/2);

business_id more visitors on weekend

Time taken: 0.065 seconds, Fetched: 182 row(s)

hive> select * from difference;

OK

05yzn45mHlTCbT25KmFVlQ	4.0
0Fln0xZBGfxvRd66pbZ3ag	8.0
0VJlqjDftGcrtgRhbkRtQA	1.0
0WsCtzG8xazAmFA5HhQE0A	3.0
1RV1FeZKJN7JJlnZpEIVWw	1.0
1Xi4hGRzhXZleq37hdFLcg	2.0
243PQVgkx8eUzVaKCgC-gw	9.0
2pm0I_mrn1ZWB88Iei0MkA	2.0
2wfiUn6AZm2K00FJYakVQ	5.0
3XBN60tLWqr7WZuDZsyZkA	6.0
49WY_TeA9ZEcRk_GnuLog	9.0
5A2BCDVbWGM7JEUgSfFleA	4.0
6FkCrxJMq5KIgavHRqtlbQ	3.0
6gDRAYnqwHDirfA5gFlsYg	2.0
6iq7LAJRSstbaYcLzQ2EzA	4.0
6Rxt0_MxnLwZ7D7iYDE8Jg	3.0
7rHrh5QggFBXj40jF0NDCA	2.0
7sakqP2av046jujj-ERHhg	8.0
Al-oKob7eDmt-_Te9FHuzw	1.0
b2596IQibXyco49ctXEupA	2.0
ba2xyZXG8hRb3hh7uDZs_g	3.0
bANAMVwSMR-UAA54WdRsIA	1.0
beQdq_Fsb-kL2e2ch7Bt6g	1.0
BGeJhokuC2wqsSKtFz47uQ	1.0
BH0K6SxPnY3wTRB7tA27hQ	2.0
b0JrPR_G0eDzpw8rGEi2mA	1.0
cB04zZj6HyWA7bmhA5Zkcg	1.0
d24LIT55-aAr7l1X99vpjw	2.0
D2Ro0b4UkH2azf8DgvNfUA	3.0
dPxL6P_Gp74CqSPuL59Srg	2.0
FpW25osFTTrQf2dCIzAPavQ	3.0
Fu-egVPTQZ3lNYM0PhpQiQ	1.0
FZ_edvdCN64D7JVWvm7fmw	3.0
gCfRc0wBLJEF1C6ilqhFuA	1.0
gHX6fK407NwP03aH9qT2AQ	2.0
glY_PSkxyM54izTHkd2mtw	1.0
hCxM2MH3cDaqozUDpydXZg	1.0
HRIUU2Z4EsZBo59P0I2ZYA	2.0
i4y0CrKYLh_Qiv6JGyABJw	1.0

STEP3:

```
sqoop import --connect jdbc:mysql://quickstart.cloudera:3306/yelp_db --username root --password cloudera --split-by business_id --table ckin --target-dir /user/cloudera/ckin --fields-terminated-by "," --hive-import --create-hive-table --hive-table yelp_db.ckin
```



```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart.cloudera:3306/yelp_db --username root --password cloudera --split-by business_id --table ckin --target-dir /user/cloudera/ckin --fields-terminated-by "," --hive-import --create-hive-table --hive-table yelp_db.ckin  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
18/05/23 08:28:11 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.12.0  
18/05/23 08:28:11 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
18/05/23 08:28:11 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
18/05/23 08:28:11 INFO tool.CodeGenTool: Beginning code generation  
18/05/23 08:28:12 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `ckin` AS t LIMIT 1  
18/05/23 08:28:12 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `ckin` AS t LIMIT 1  
18/05/23 08:28:12 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce  
Note: /tmp/sqoop-cloudera/compile/01d3dcdb2f470ccdf5a1b21470bc4b9/ckin.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
18/05/23 08:28:14 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/01d3dcdb2f470ccdf5a1b21470bc4b9/ckin.jar  
18/05/23 08:28:14 WARN manager.MySQLManager: It looks like you are importing from mysql.  
18/05/23 08:28:14 WARN manager.MySQLManager: This transfer can be faster! Use the --direct  
18/05/23 08:28:14 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.  
18/05/23 08:28:14 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)  
18/05/23 08:28:14 INFO mapreduce.ImportJobBase: Beginning import of ckin  
18/05/23 08:28:14 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address  
18/05/23 08:28:14 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
18/05/23 08:28:15 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
18/05/23 08:28:15 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032  
18/05/23 08:28:17 INFO db.DBInputFormat: Using read committed transaction isolation  
18/05/23 08:28:17 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`business_id`), MAX(`business_id`) FROM `ckin`  
18/05/23 08:28:17 WARN db.TextSplitter: Generating splits for a textual index column.  
18/05/23 08:28:17 WARN db.TextSplitter: If your database sorts in a case-insensitive order, this may result in a partial import or duplicate records.  
18/05/23 08:28:17 WARN db.TextSplitter: You are strongly encouraged to choose an integral split column.  
18/05/23 08:28:17 INFO mapreduce.JobSubmitter: number of splits:6  
18/05/23 08:28:17 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1526588254685_0004  
18/05/23 08:28:18 INFO impl.YarnClientImpl: Submitted application application_1526588254685_0004  
18/05/23 08:28:18 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1526588254685_0004/  
18/05/23 08:28:18 INFO mapreduce.Job: Running job: job_1526588254685_0004  
18/05/23 08:28:26 INFO mapreduce.Job: Job job_1526588254685_0004 running in uber mode : false  
18/05/23 08:28:26 INFO mapreduce.Job: map 0% reduce 0%  
18/05/23 08:28:44 INFO mapreduce.Job: map 17% reduce 0%  
18/05/23 08:28:50 INFO mapreduce.Job: map 33% reduce 0%  
18/05/23 08:28:52 INFO mapreduce.Job: map 50% reduce 0%  
18/05/23 08:28:54 INFO mapreduce.Job: map 100% reduce 0%  
18/05/23 08:28:54 INFO mapreduce.Job: Job job_1526588254685_0004 completed successfully  
18/05/23 08:28:54 INFO mapreduce.Job: Counters: 31  
File System Counters
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
FILE: Number of large read operations=0  
FILE: Number of write operations=0  
HDFS: Number of bytes read=857  
HDFS: Number of bytes written=109593  
HDFS: Number of read operations=24  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=12  
Job Counters  
  Killed map tasks=1  
  Launched map tasks=6  
  Other local map tasks=6  
  Total time spent by all maps in occupied slots (ms)=127955  
  Total time spent by all reduces in occupied slots (ms)=0  
  Total time spent by all map tasks (ms)=127955  
  Total vcore-milliseconds taken by all map tasks=127955  
  Total megabyte-milliseconds taken by all map tasks=131025920  
Map-Reduce Framework  
  Map input records=3257  
  Map output records=3257  
  Input split bytes=857  
  Spilled Records=0  
  Failed Shuffles=0  
  Merged Map outputs=0  
  GC time elapsed (ms)=2397  
  CPU time spent (ms)=4760  
  Physical memory (bytes) snapshot=958472192  
  Virtual memory (bytes) snapshot=9061216256  
  Total committed heap usage (bytes)=748683264  
File Input Format Counters  
  Bytes Read=0  
File Output Format Counters  
  Bytes Written=109593  
18/05/23 08:28:54 INFO mapreduce.ImportJobBase: Transferred 107.0244 KB in 39.0609 seconds (2.7399 KB/sec)  
18/05/23 08:28:54 INFO mapreduce.ImportJobBase: Retrieved 3257 records.  
18/05/23 08:28:54 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `ckin` AS t LIMIT 1  
18/05/23 08:28:54 INFO hive.HiveImport: Loading uploaded data into Hive  
  
Logging initialized using configuration in jar:file:/usr/lib/hive/lib/hive-common-1.1.0-cdh5.12.0.jar!/hive-log4j.properties  
OK  
Time taken: 1.959 seconds  
Loading data to table yelp.db.ckin  
Table yelp_db.ckin stats: [numFiles=6, totalSize=109593]  
OK
```

Step4:

```
sqoop import --connect jdbc:mysql://quickstart.cloudera:3306/yelp_db --username root --password cloudera --split-by business_id --table ckin_pivot --target-dir /user/cloudera/ckin_pivot --fields-terminated-by "," --hive-import --create-hive-table --hive-table yelp_db.ckin_pivot
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
[cloudera@quickstart ~]$ sqoop import --connect jdbc:mysql://quickstart.cloudera:3306/yelp_db --username root --password cloudera --split-by business_id --table ckin_pivot --get-dir /user/cloudera/ckin_pivot --fields-terminated-by "," --hive-import --create-hive-table --hive-table yelp_db.ckin_pivot  
Warning: /usr/lib/sqoop/./accumulo does not exist! Accumulo imports will fail.  
Please set $ACCUMULO_HOME to the root of your Accumulo installation.  
18/05/23 08:48:38 INFO sqoop.Sqoop: Running Sqoop version: 1.4.6-cdh5.12.0  
18/05/23 08:48:38 WARN tool.BaseSqoopTool: Setting your password on the command-line is insecure. Consider using -P instead.  
18/05/23 08:48:39 INFO manager.MySQLManager: Preparing to use a MySQL streaming resultset.  
18/05/23 08:48:39 INFO tool.CodeGenTool: Beginning code generation  
18/05/23 08:48:39 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `ckin_pivot` AS t LIMIT 1  
18/05/23 08:48:39 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `ckin_pivot` AS t LIMIT 1  
18/05/23 08:48:39 INFO orm.CompilationManager: HADOOP_MAPRED_HOME is /usr/lib/hadoop-mapreduce  
Note: /tmp/sqoop-cloudera/compile/3940a3258e622cce2d6ff6b87c9a6a4f/ckin_pivot.java uses or overrides a deprecated API.  
Note: Recompile with -Xlint:deprecation for details.  
18/05/23 08:48:42 INFO orm.CompilationManager: Writing jar file: /tmp/sqoop-cloudera/compile/3940a3258e622cce2d6ff6b87c9a6a4f/ckin_pivot.jar  
18/05/23 08:48:42 WARN manager.MySQLManager: It looks like you are importing from mysql.  
18/05/23 08:48:42 WARN manager.MySQLManager: This transfer can be faster! Use the --direct  
18/05/23 08:48:42 WARN manager.MySQLManager: option to exercise a MySQL-specific fast path.  
18/05/23 08:48:42 INFO manager.MySQLManager: Setting zero DATETIME behavior to convertToNull (mysql)  
18/05/23 08:48:42 INFO mapreduce.ImportJobBase: Beginning import of ckin_pivot  
18/05/23 08:48:42 INFO Configuration.deprecation: mapred.job.tracker is deprecated. Instead, use mapreduce.jobtracker.address  
18/05/23 08:48:42 INFO Configuration.deprecation: mapred.jar is deprecated. Instead, use mapreduce.job.jar  
18/05/23 08:48:43 INFO Configuration.deprecation: mapred.map.tasks is deprecated. Instead, use mapreduce.job.maps  
18/05/23 08:48:43 INFO client.RMProxy: Connecting to ResourceManager at /0.0.0.0:8032  
18/05/23 08:48:44 WARN hdfs.DFSClient: Caught exception  
java.lang.InterruptedException  
    at java.lang.Object.wait(Native Method)  
    at java.lang.Thread.join(Thread.java:1281)  
    at java.lang.Thread.join(Thread.java:1355)  
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.closeResponder(DFSOutputStream.java:952)  
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.endBlock(DFSOutputStream.java:690)  
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:879)  
18/05/23 08:48:45 INFO db.DBInputFormat: Using read committed transaction isolation  
18/05/23 08:48:45 INFO db.DataDrivenDBInputFormat: BoundingValsQuery: SELECT MIN(`business_id`), MAX(`business_id`) FROM `ckin_pivot`  
18/05/23 08:48:45 WARN db.TextSplitter: Generating splits for a textual index column.  
18/05/23 08:48:45 WARN db.TextSplitter: If your database sorts in a case-insensitive order, this may result in a partial import or duplicate records.  
18/05/23 08:48:45 WARN db.TextSplitter: You are strongly encouraged to choose an integral split column.  
18/05/23 08:48:45 INFO mapreduce.JobSubmitter: number of splits:6  
18/05/23 08:48:45 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1526588254685_0005  
18/05/23 08:48:45 INFO impl.YarnClientImpl: Submitted application application_1526588254685_0005  
18/05/23 08:48:45 INFO mapreduce.Job: The url to track the job: http://quickstart.cloudera:8088/proxy/application_1526588254685_0005/  
18/05/23 08:48:45 INFO mapreduce.Job: Running job: job_1526588254685_0005  
18/05/23 08:48:53 INFO mapreduce.Job: Job job_1526588254685_0005 running in uber mode : false  
18/05/23 08:48:53 INFO mapreduce.Job: map 0% reduce 0%
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
HDFS: Number of bytes read=857  
HDFS: Number of bytes written=6801  
HDFS: Number of read operations=24  
HDFS: Number of large read operations=0  
HDFS: Number of write operations=12  
Job Counters  
Killed map tasks=1  
Launched map tasks=6  
Other local map tasks=6  
Total time spent by all maps in occupied slots (ms)=122785  
Total time spent by all reduces in occupied slots (ms)=0  
Total time spent by all map tasks (ms)=122785  
Total vcore-milliseconds taken by all map tasks=122785  
Total megabyte-milliseconds taken by all map tasks=125731840  
Map-Reduce Framework  
Map input records=182  
Map output records=182  
Input split bytes=857  
Spilled Records=0  
Failed Shuffles=0  
Merged Map outputs=0  
GC time elapsed (ms)=2031  
CPU time spent (ms)=4470  
Physical memory (bytes) snapshot=949485568  
Virtual memory (bytes) snapshot=9060782080  
Total committed heap usage (bytes)=748683264  
File Input Format Counters  
Bytes Read=0  
File Output Format Counters  
Bytes Written=6801  
18/05/23 08:49:21 INFO mapreduce.ImportJobBase: Transferred 6.6416 KB in 38.1557 seconds (178.2436 bytes/sec)  
18/05/23 08:49:21 INFO mapreduce.ImportJobBase: Retrieved 182 records.  
18/05/23 08:49:21 INFO manager.SqlManager: Executing SQL statement: SELECT t.* FROM `ckin_pivot` AS t LIMIT 1  
18/05/23 08:49:21 INFO hive.HiveImport: Loading uploaded data into Hive  
  
Logging initialized using configuration in jar:file:/usr/lib/hive/lib/hive-common-1.1.0-cdh5.12.0.jar!/hive-log4j.properties  
OK  
Time taken: 1.76 seconds  
Loading data to table yelp_db.ckin_pivot  
Table yelp_db.ckin_pivot stats: [numFiles=6, totalSize=6801]  
OK  
Time taken: 0.386 seconds  
[cloudera@quickstart ~]$
```

```
cloudera@quickstart:~  
File Edit View Search Terminal Help  
hive> SET hive.auto.convert.join=false;  
hive> select b.name,b.city,b.state,d.weekendmorevisitors as visitors_on_weekend from difference d,business b where d.business_id = b.id;  
Query ID = cloudera_20180523092525_82879134-8236-4a3e-8006-a8687180b827  
Total jobs = 1  
Launching Job 1 out of 1  
Number of reduce tasks not specified. Estimated from input data size: 1  
In order to change the average load for a reducer (in bytes):  
    set hive.exec.reducers.bytes.per.reducer=<number>  
In order to limit the maximum number of reducers:  
    set hive.exec.reducers.max=<number>  
In order to set a constant number of reducers:  
    set mapreduce.job.reduces=<number>  
Starting Job = job_1526588254685_0008, Tracking URL = http://quickstart.cloudera:8088/proxy/application_1526588254685_0008/  
Kill Command = /usr/lib/hadoop/bin/hadoop job -kill job_1526588254685_0008  
Hadoop job information for Stage-1: number of mappers: 2; number of reducers: 1  
2018-05-23 09:25:08,614 Stage-1 map = 0%, reduce = 0%  
2018-05-23 09:25:19,572 Stage-1 map = 50%, reduce = 0%, Cumulative CPU 1.22 sec  
2018-05-23 09:25:20,605 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 2.33 sec  
2018-05-23 09:25:28,084 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 3.84 sec  
MapReduce Total cumulative CPU time: 3 seconds 840 msec  
Ended Job = job_1526588254685_0008  
MapReduce Jobs Launched:  
Stage-Stage-1: Map: 2 Reduce: 1 Cumulative CPU: 3.84 sec HDFS Read: 70209 HDFS Write: 3334 SUCCESS  
Total MapReduce CPU Time Spent: 3 seconds 840 msec  
OK  
Game Seven Grill      Phoenix AZ      4.0  
Circle K              Phoenix AZ      8.0  
Smile! Dental Center  Ste 30"        Gibsonia        1.0  
Happy Moose Bar and Grill Streetsboro OH    3.0  
Beaumont Tire & Auto Repair Champaign IL     1.0  
Jody Maroni's Sausage Kingdom 3790 Las Vegas Blvd S" Las Vegas 2.0  
Travelodge            Scottsdale AZ    9.0  
Tim Hortons           Markham ON       2.0  
Men's Ultimate Grooming Ste 103" Gilbert 5.0  
Epic Thrift           Las Vegas NV     6.0  
Politic Chester CHE   9.0  
Double Play Sports Bar Las Vegas NV     4.0  
The Independent 1 Floor" Toronto 3.0  
Knitting Tree         Madison WI       3.0  
Rôtisserie Benny & Co Sainte-Julie QC   2.0  
Wong's Jr            Peoria AZ       4.0  
Rite Aid             Pittsburgh PA    2.0  
Travel Leaders       North Olmsted OH 8.0
```

STEP5:

OUTPUT:

select b.name, b.city,b.state, d.weekendmorevisitors as visitors_on_weekend from difference d,business b where d.business_id = b.id;

I have executed sqoop jobs to import data from my sql to hive database. The Map reduce jobs are executed

name	city	state	visitors_on_weekend
Game Seven Grill	Phoenix	AZ	4
Circle K	Phoenix	AZ	8
Smile! Dental Center	Gibsonia	PA	1
Happy Moose Bar and Grill	Streetsboro	OH	3
Beaumont Tire & Auto Repair	Champaign	IL	1
Jody Maroni's Sausage Kingdom	Las Vegas	NV	2
Travelodge	Scottsdale	AZ	9
Tim Hortons	Markham	ON	2
Men's Ultimate Grooming	Gilbert	AZ	5
Epic Thrift	Las Vegas	NV	6
Politic	Chester	CHE	9
Double Play Sports Bar	Las Vegas	NV	4
The Independent	Toronto	ON	3
Rôtisserie Benny & Co	Sainte-Julie	QC	2
Wong's Jr	Peoria	AZ	4
Knitting Tree	Madison	WI	3
Rite Aid	Pittsburgh	PA	2
Travel Leaders	North Olmsted	OH	8
Bridal Fashion Debut	Phoenix	AZ	1
URBAN JUNGLE FUN PARK Mesa	Mesa	AZ	2
Mike Olbinski Photography	Phoenix	AZ	3
Advantage Rent-A-Car	Phoenix	AZ	1
Deli 194	Edinburgh	EDH	1
Boulangerie Arouch	Laval	QC	1
Sunnyside Grill	Toronto	ON	2
The Fry	North York	ON	1
Centra Food Market	Aurora	ON	1
Domino's Pizza	Tempe	AZ	2
New Wrap Order	Charlotte	NC	3
World Food Championships	Las Vegas	NV	2
Catholic Charities	Las Vegas	NV	3
MV Nail Spa	North Olmsted	OH	1
Jose' and Tony's Mexican Restaurant	Pittsburgh	PA	3
Mulberry Bar	Toronto	ON	1
Wendy's	Stow	OH	2
Matt's Big Breakfast	Phoenix	AZ	1
Zonie Cuts	Phoenix	AZ	1